

B211038- Srushti Gavale

Boston house price

```
# Importing libraries
import numpy as np
import pandas as pd
import tensorflow as tf
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import r2_score

# Loading the Boston Housing dataset
boston_dataset = pd.read_csv('Boston.csv')
boston = pd.DataFrame(boston_dataset, columns=boston_dataset.columns)
boston['MEDV'] = boston_dataset['medv']

boston_dataset.shape

(506, 16)

print(boston_dataset.head(5))
```

	Unnamed: 0	crim	zn	indus	chas	nox	rm	age	dis	rad	\
0	1	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	
1	2	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	
2	3	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	
3	4	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	
4	5	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	

```

tax  ptratio  black  lstat  medv  MEDV
0  296      15.3  396.90  4.98  24.0  24.0
1  242      17.8  396.90  9.14  21.6  21.6
2  242      17.8  392.83  4.03  34.7  34.7
3  222      18.7  394.63  2.94  33.4  33.4
4  222      18.7  396.90  5.33  36.2  36.2

print(np.shape(boston_dataset))

(506, 16)

print(boston_dataset.describe())
```

	Unnamed: 0	crim	zn	indus	chas	nox	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	253.500000	3.613524	11.363636	11.136779	0.069170	0.554695	
std	146.213884	8.601545	23.322453	6.860353	0.253994	0.115878	
min	1.000000	0.006320	0.000000	0.460000	0.000000	0.385000	
25%	127.250000	0.082045	0.000000	5.190000	0.000000	0.449000	
50%	253.500000	0.256510	0.000000	9.690000	0.000000	0.538000	
75%	379.750000	3.677083	12.500000	18.100000	0.000000	0.624000	
max	506.000000	88.976200	100.000000	27.740000	1.000000	0.871000	

```

rm      age      dis      rad      tax      ptratio  \
count  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000
mean    6.284634    68.574901    3.795043    9.549407    408.237154    18.455534
std     0.702617    28.148861    2.105710    8.707259    168.537116    2.164946
min     3.561000    2.900000    1.129600    1.000000    187.000000    12.600000
25%     5.885500    45.025000    2.100175    4.000000    279.000000    17.400000
50%     6.208500    77.500000    3.207450    5.000000    330.000000    19.050000
75%     6.623500    94.075000    5.188425    24.000000    666.000000    20.200000
max     8.780000   100.000000   12.126500   24.000000   711.000000   22.000000

black    lstat    medv    MEDV
count  506.000000  506.000000  506.000000  506.000000
mean   356.674032  12.653063   22.532806  22.532806
std    91.294864   7.141062   9.197104   9.197104
min     0.320000   1.730000   5.000000   5.000000
25%    375.377500   6.950000  17.025000  17.025000
50%    391.440000  11.360000  21.200000  21.200000
75%    396.225000  16.955000  25.000000  25.000000
max    396.900000  37.970000  50.000000  50.000000

# Split the data into training and testing sets
X = boston.drop('MEDV', axis=1)
Y = boston['MEDV']
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=1)
```

```

# Scale the data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Define the model
model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(64, activation='relu', input_shape=(X_train.shape[1],)),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(1)
])

# Compile the model
model.compile(optimizer='adam', loss='mse')

# Train the model
history = model.fit(X_train_scaled, Y_train, validation_data=(X_test_scaled, Y_test), epochs=100)

Epoch 1/100
13/13 [=====] - 2s 22ms/step - loss: 579.7935 - val_loss: 566.8345
Epoch 2/100
13/13 [=====] - 0s 4ms/step - loss: 519.9058 - val_loss: 499.8945
Epoch 3/100
13/13 [=====] - 0s 4ms/step - loss: 446.0014 - val_loss: 410.6792
Epoch 4/100
13/13 [=====] - 0s 4ms/step - loss: 346.0971 - val_loss: 295.0309
Epoch 5/100
13/13 [=====] - 0s 4ms/step - loss: 228.5537 - val_loss: 172.4942
Epoch 6/100
13/13 [=====] - 0s 4ms/step - loss: 117.2006 - val_loss: 86.2677
Epoch 7/100
13/13 [=====] - 0s 4ms/step - loss: 60.7959 - val_loss: 50.8605
Epoch 8/100
13/13 [=====] - 0s 4ms/step - loss: 41.9353 - val_loss: 37.0392
Epoch 9/100
13/13 [=====] - 0s 4ms/step - loss: 31.4906 - val_loss: 27.4767
Epoch 10/100
13/13 [=====] - 0s 3ms/step - loss: 24.2087 - val_loss: 21.7079
Epoch 11/100
13/13 [=====] - 0s 4ms/step - loss: 19.9854 - val_loss: 18.1820
Epoch 12/100
13/13 [=====] - 0s 4ms/step - loss: 17.1946 - val_loss: 15.7325
Epoch 13/100
13/13 [=====] - 0s 4ms/step - loss: 15.0696 - val_loss: 14.0107
Epoch 14/100
13/13 [=====] - 0s 4ms/step - loss: 13.4463 - val_loss: 12.7971
Epoch 15/100
13/13 [=====] - 0s 4ms/step - loss: 12.3516 - val_loss: 11.7113
Epoch 16/100
13/13 [=====] - 0s 4ms/step - loss: 11.1735 - val_loss: 10.7977
Epoch 17/100
13/13 [=====] - 0s 4ms/step - loss: 10.3415 - val_loss: 10.2638
Epoch 18/100
13/13 [=====] - 0s 4ms/step - loss: 9.5420 - val_loss: 9.5540
Epoch 19/100
13/13 [=====] - 0s 4ms/step - loss: 8.9082 - val_loss: 9.0734
Epoch 20/100
13/13 [=====] - 0s 4ms/step - loss: 8.3447 - val_loss: 8.7130
Epoch 21/100
13/13 [=====] - 0s 4ms/step - loss: 7.7903 - val_loss: 8.1594
Epoch 22/100
13/13 [=====] - 0s 4ms/step - loss: 7.3412 - val_loss: 7.7143
Epoch 23/100
13/13 [=====] - 0s 4ms/step - loss: 6.9347 - val_loss: 7.3838
Epoch 24/100
13/13 [=====] - 0s 4ms/step - loss: 6.5723 - val_loss: 7.0898
Epoch 25/100
13/13 [=====] - 0s 4ms/step - loss: 6.2431 - val_loss: 6.6430
Epoch 26/100
13/13 [=====] - 0s 4ms/step - loss: 5.9225 - val_loss: 6.4520
Epoch 27/100
13/13 [=====] - 0s 4ms/step - loss: 5.6786 - val_loss: 6.1964
Epoch 28/100
13/13 [=====] - 0s 4ms/step - loss: 5.3824 - val_loss: 5.9290
Epoch 29/100
13/13 [=====] - 0s 3ms/step - loss: 5.1993 - val_loss: 5.7668

# Evaluate the model
Y_pred = model.predict(X_test_scaled)
r2 = r2_score(Y_test, Y_pred)
print("R^2 score:", r2)

```

```
4/4 [=====] - 0s 2ms/step
R^2 score: 0.987125595222069
```

```
import numpy as np
import seaborn as sns

# Generate some sample data
X = np.random.normal(0, 1, 100)
Y = 2 * X + np.random.normal(0, 1, 100)

# Fit a linear regression model
model = np.polyfit(X, Y, 1)

# Make predictions on the training data
Y_pred = np.polyval(model, X)

# Add axis labels
plt.xlabel('True Values')
plt.ylabel('Predicted Values')

# Create a scatter plot of predicted vs true values
sns.scatterplot(np.squeeze(Y), np.squeeze(Y_pred))

# Add a diagonal line to show perfect correlation
sns.lineplot(np.squeeze(Y), np.squeeze(Y), color='red')
```

```
C:\Users\D_COMP_RSL-14\anaconda3\lib\site-packages\seaborn\_decorators.py:36: FutureWarning: Pass the following variables a
warnings.warn(
C:\Users\D_COMP_RSL-14\anaconda3\lib\site-packages\seaborn\_decorators.py:36: FutureWarning: Pass the following variables a
warnings.warn(
<AxesSubplot: xlabel='True Values', ylabel='Predicted Values'>
```



